


Lec. 1: Shortest Path & Min-Mean Circuit



Walk, Path, Circuits

- Walk: a traversal from vertex to vertex along edges.
- Trail: a walk without repeated edges
- Path: a walk without repeated internal vertices
- Cycle: closed trail
- Circuit: closed path




Shortest Paths

Shortest path problem. Given a directed graph $D = (V, A)$ with edge lengths (costs, weights) $\ell(u, v)$, find shortest path from node s to node t .

↖ allow negative weights

Ex. Nodes represent agents in a financial setting and $\ell(u, v)$ is cost of transaction in which we buy from agent u and sell immediately to v .



NP-Completeness of Shortest Path

NP-complete: even if each arc has length -1 . Equivalently, finding a longest path in a graph (with unit length arcs) is NP-complete.

Pf. Reduction from finding a Hamiltonian path

Remark: A shortest **walk** with at most (resp. exactly) k arcs can be computed in polynomial time.

NP-Completeness of Shortest Circuit

NP-complete: even if each arc has length -1. Equivalently, finding a Hamiltonian circuit in a graph (with unit length arcs) is NP-complete.

Minimum-Mean Circuit: a circuit C with the least mean length $\ell(C)/|C|$.
Solvable in polynomial time


Outline

- Shortest path with unit lengths
- Shortest path with non-negative lengths
- Shortest Walk with arbitrary lengths
- All-pairs shortest paths
- Minimum-mean length directed circuit
- Elementary decomposition of circulations and transshipments

1. Shortest Path: Unit Lengths

SP with Unit Lengths: Breadth First Search

BFS intuition. Explore outward from s in all possible directions, adding nodes one "layer" at a time.





BFS algorithm.

- $L_0 = \{ s \}$.
- L_1 = all neighbors of L_0 .
- L_2 = all nodes that do not belong to L_0 or L_1 , and that have an edge to a node in L_1 .
- L_{i+1} = all nodes that do not belong to an earlier layer, and that have an edge to a node in L_i .


Theorem. For each i , L_i consists of all nodes at distance exactly i from s .

Shortest-Path Tree


Theorem. The A BFS tree is a shortest-path tree and can be computed in $O(m + n)$ time if the graph is given by its adjacency list.



(a)



(b)



(c)

L_0


L_1

L_2

L_3

Level Graph

- Subgraph of D consisting of all vertices and edges appearing in some shortest s - t path in D
- Compute in $O(m + n)$ time using BFS by keeping only forward edges (deleting back and side edges).
- An inclusion-wise maximal collection of edge-disjoint (or vertex-disjoint) shortest s - t paths can be computed in $O(m + n)$ time (exercise).



2. Shortest-Path: Non-Negative Lengths


Dijkstra's (Greedy) Algorithm

- Maintain a set of **explored nodes** S for which we have determined the shortest path distance $d(u)$ from s to u .
- Initialize $S = \{s\}$, $d(s) = 0$.
- Repeatedly choose unexplored node v which minimizes

$$\pi(v) = \min_{e=(u,v): u \in S} d(u) + \ell(e),$$

add v to S , and set $d(v) = \pi(v)$.

← shortest path to some u in explored part, followed by a single edge (u, v)



Dijkstra's Algorithm: Push-Based Implementation

For each unexplored node, explicitly maintain $\pi(v) = \min_{e=(u,v): u \in S} d(u) + \ell(e)$.

- Next node to explore = node with minimum $\pi(v)$.
- When exploring v , for each $e = (v, w)$ update

$$\pi(w) = \min \{ \pi(w), \pi(v) + \ell(e) \}.$$

Efficient implementation. Maintain a priority queue of unexplored nodes, prioritized by $\pi(v)$.

PQ Operation	Dijkstra	Array	Binary heap	d-way Heap	Fib heap [†]
Insert	n	n	$\log n$	$d \log_d n$	1
ExtractMin	n	n	$\log n$	$d \log_d n$	$\log n$
ChangeKey	m	1	$\log n$	$\log_d n$	1
IsEmpty	n	1	1	1	1
Total		n^2	$m \log n$	$m \log_{m/n} n$	$m + n \log n$

[†] Individual ops are amortized bounds

Reweighting Edges with Node Prices

p : a node price function

p -adjusted edge length: for each edge $a = (u, v)$

$$\ell_p(a) := \ell(a) - p(v) + p(u)$$

- purchase from u at price $p(u)$, ship to v at price $\ell(a)$, and sell to v at price $p(v)$

Invariant Properties:

- For each cycle C , $\ell_p(C) = \ell(C)$
- For each s - t walk/path P , $\ell_p(P) = \ell(P) - p(t) + p(s)$

Each cycle or s - t walk/path is shortest w.r.t. $\ell_p \iff$
it is shortest w.r.t. ℓ .

Node Potential

Wish: ℓ_p is nonnegative s.t. Dijkstra's algorithm can be applied.

Potential p : $p(v) - p(u) \leq \ell(a)$ for each arc $a = (u, v)$;
 $\Leftrightarrow \ell_p$ is nonnegative.


Theorem. There exists a potential \Leftrightarrow each circuit is nonnegative.
If moreover ℓ is integer, the potential can be taken integer.

Pf. \Leftarrow distance-based potential: $p(v) :=$ the s-v distance.
 \Rightarrow trivial


Remark: distance-based potential maximizes $p(t) - p(s)$ (string model)

3. Shortest Walks: Arbitrary Lengths

Negative Cycles



Observation. If some s-t **walk** contains a negative cycle, there does not exist a shortest s-t **walk**; otherwise, there exists one that is a s-t **path**.



Bellman-Ford Algorithm (DP)

Def. $d_i(v)$ = length of shortest v - t walk P using at most i edges.

$$d_0(v) = \begin{cases} 0 & \text{if } v = t \\ \infty & \text{otherwise} \end{cases}$$

For $i > 0$

$$d_i(v) = \min \left\{ d_{i-1}(v), \min_{(v,w) \in A} \{ d_{i-1}(w) + \ell(vw) \} \right\}$$

Remark. By previous observation, if no negative cycles, then $d_{n-1}(v)$ = length of shortest v - t path.

Naïve Implementation

```
Shortest-Walk(D, t) {  
  foreach node v ∈ V  
    d[0, v] ← ∞  
  d[0, t] ← 0  
  
  for i = 1 to n-1  
    foreach node v ∈ V  
      d[i, v] ← d[i-1, v]  
    foreach edge (v, w) ∈ A  
      d[i, v] ← min { d[i, v], d[i-1, w] + ℓ(vw) }  
}
```

Analysis. $\Theta(mn)$ time, $\Theta(n^2)$ space.

Finding the shortest walks. Maintain a "successor" for each table entry.

Space-Efficient Improvement with 1D-Table

- Maintain only 1D array $d[v]$ = shortest v - t walk found so far.
- No need to check (v, w) unless $d[w]$ changed in previous iteration.

```
Push-Based-Shortest-Walk( $D, s, t$ ) {  
  foreach node  $v \in V$   
     $d[v] \leftarrow \infty$ ,  $\text{successor}[v] \leftarrow \phi$   
   $d[t] = 0$   
  
  for  $i = 1$  to  $n-1$   
    foreach node  $w \in V$   
      if ( $d[w]$  has been updated in previous iteration)  
        foreach node  $v$  such that  $(v, w) \in A$   
          if ( $d[v] > d[w] + \ell(vw)$ )  
             $d[v] \leftarrow d[w] + \ell(vw)$ ,  $\text{successor}[v] \leftarrow w$   
  If no  $d[w]$  value changed in iteration  $i$ , stop.
```

Space-Efficient Improvement with 1D-Table

Theorem. Throughout the algorithm, each finite $d[v]$ is length of some v - t walk; and after i rounds of updates, $d[v] \leq d_i(v)$.

Pf. By induction on i .

Overall impact.

- Memory: $O(m + n)$.
- Running time: $O(mn)$ worst case, but substantially faster in practice

Detecting Negative Cycles


Lemma. If $d_n(v) = d_{n-1}(v)$ for all v , then no negative cycles.

Pf. By Bellman-Ford algorithm, $d_i(v) = d_{n-1}(v)$ for all $i \geq n - 1$.

Lemma. If $d_n(v) < d_{n-1}(v)$ for some node v , then (any) shortest walk of n arcs from v to t contains a cycle C . Moreover, C has negative cost.

Pf. (by contradiction)


- Since $d_n(v) < d_{n-1}(v)$, we know P has exactly n edges.
- By pigeonhole principle, P must contain a directed cycle C .
- Deleting C yields a v - t path with $< n$ edges $\Rightarrow C$ has negative cost.



Detecting Negative Cycles

Detect negative cycle in $O(mn)$ time.

- Add new node t and connect all nodes to t with 0-length edge.
- Check if $d_n(v) = d_{n-1}(v)$ for all nodes v .
 - if yes, then no negative cycles
 - if no, then extract cycle from shortest walk from v to t



Detecting Negative Cycles: Summary

Bellman-Ford. $O(mn)$ time, $O(m + n)$ space.

- Run Bellman-Ford for n iterations (instead of $n-1$).
- Upon termination, Bellman-Ford successor variables trace a negative cycle if one exists.

Shortest k-Link Walks (DP)

Def. $d_i(v)$ = minimum length of a walk ending at v with exactly i arcs

$$d_i(v) = \begin{cases} 0 & \text{if } i = 0 \\ \min_{(u,v) \in A} \{d_{i-1}(u) + \ell(uv)\} & \text{otherwise} \end{cases}$$

```
foreach v ∈ V
    d[0,v] ← 0
    predecessor[0,v] ← φ

for k = 1 to n
    foreach v ∈ V
        d[k,v] ← ∞
        foreach (u,v) ∈ A
            if d[k,v] > d[k-1,v] + ℓ(uv) then
                d[k,v] ← d[k-1,v] + ℓ(uv)
                predecessor[k,v] ← u
```

4. All-Pairs Shortest Paths

Floyd-Warshall method (DP)

v_1, v_2, \dots, v_n : an arbitrary vertex ordering

$d_k(s, t) :=$ minimum length of an s - t **walk** using only vertices in $\{s, t, v_1, \dots, v_k\}$.

$$d_0(s, t) = \begin{cases} \ell(s, t) & \text{if } (s, t) \in A \\ \infty & \text{otherwise} \end{cases}$$

$$d_k(s, t) = \min \{d_{k-1}(s, t), d_{k-1}(s, v_k) + d_{k-1}(v_k, t)\}$$

$$dist_\ell = d_n$$

Theorem. Under the condition of no negative-length circuit, all distances can be determined in time $O(n^3)$.

Space-Efficient Implementation with 2D-Table

FW-Shortest-Path(D)

```
foreach (u,v) ∈ V × V
  if (u,v) ∈ A then
    d[u,v] ← ℓ(u,v), successor[u,v] ← v
  else
    d[u,v] ← ∞, successor[u,v] ← ∅

for k = 1 to n
  for i = 1 to n
    for j = 1 to n
      if d[i,j] > d[i,k] + d[k,j] then
        d[i,j] ← d[i,k] + d[k,j]
        successor[i,j] ← successor[i,k]
```

Theorem. Throughout the algorithm, each finite $d[u,v]$ is length of some u - v walk; and after k rounds of updates, $d[u,v] \leq d_k(u,v)$.

Johnson's Algorithm (DP+Greedy)

Johnson's algorithm.

- Apply the Bellman-Ford method to find a distance-based potential p
- Reweight the lengths with p
- Apply Dijkstra's method to compute a shortest-path tree rooted at each other node

Analysis. $\Theta(n(m + n \log n))$ time.

5. Minimum-Mean Circuit

Minimum-Mean Circuit (MMC)

Def. mean length of a circuit $C := \ell(C)/|C|$.

Min-mean = the smallest value such that can be subtracted from each edge to ensure that each circuit becomes nonnegative.

MMC: invariant under **uniform** change on edge length

Assumption: D is strongly connected, for otherwise consider individual strong components.

Minimum Mean Length

Def. $d_i(v)$ = minimum length of a **walk** ending at v with **exactly** i arcs


Theorem. The minimum mean length of a circuit is equal to

$$\min_{v \in V} \max_{0 \leq i \leq n-1} \frac{d_n(v) - d_i(v)}{n - i}.$$

Pf. W.l.o.g. assume MMC length = 0 (hence minimum circuit length = 0).

$$\min_{v \in V} [d_n(v) - \min_{0 \leq i \leq n-1} d_i(v)] = 0.$$


□ For **all** v , $d_n(v) \geq \min_{0 \leq i \leq n-1} d_i(v)$



Minimum Mean Length

Pf. For **some** v , $d_n(v) \leq \min_{0 \leq i \leq n-1} d_n(v)$

- $\min_j d_j(v_0)$ is attained by some **j** with $n-k \leq j < n$.
- $\min_i d_i(v_{n-j})$ is attained by some i with $0 \leq i < n$.



$$d_n(v_{n-j}) \leq d_j(v_0) + \ell(P) \leq d_i(v_{n-j}) + \ell(Q) + \ell(P) = d_i(v_{n-j})$$

Karp's Algorithm

- Compute $d_i(v)$ for all v and $0 \leq i \leq n$.
- Find v minimizing $\max_{0 \leq i \leq n-1} \frac{d_n(v) - d_i(v)}{n-i}$
 - Compute a shortest walk P ending at v with **exactly** n arcs,
 - Find a circuit C in P and output C

$P-C$ is a walk ending at v with $i := n - |C|$ arcs

$$\frac{\ell(C)}{|C|} = \frac{\ell(P) - \ell(P-C)}{n-i} \leq \frac{d_n(v) - d_i(v)}{n-i}$$

Running time: $O(mn)$


6. Decomposition of Circulations

Non-negative Circulation

$D = (V, A)$: a simple directed graph

x : a **non-negative** edge function


Def. x is a **circulation** if for any $v \in V$, $x(\delta^{in}(v)) = x(\delta^{out}(v))$



Elementary Decomposition

Elementary circulation: χ_C along a circuit C in A

- $\chi_C(a) = 1$ for each $a \in C$;
- $\chi_C(a) = 0$ for each $a \notin C$.



$$A^+(x) := \{a \in A : x(a) > 0\}$$


Theorem: There exist $k \leq |A^+(x)|$ **circuits** C_1, \dots, C_k in $A^+(x)$ together with k positive numbers $\varepsilon_1, \dots, \varepsilon_k$ s.t.

$$x = \sum_{i=1}^k \varepsilon_i \chi_{C_i}.$$


Moreover, if x is integer-valued, so are all ε -values.

The decomposition can be found in $O(nm)$ time.


Elementary Decomposition of Circulation




Elementary Decomposition of Circulation




Elementary Decomposition of Circulation



Elementary Decomposition of Circulation




Elementary Decomposition of Circulation




1,2,5,8 : 9
1,4,7,8 : 10

Elementary Decomposition of Circulation




1,2,5,8 : 9
1,4,7,8 : 10

Elementary Decomposition of Circulation




1,2,5,8 : 9
1,4,7,8 : 10
1,3,6,8 : 4

Elementary Decomposition of Circulation




1,2,5,8 : 9
1,4,7,8 : 10
1,3,6,8 : 4

Elementary Decomposition of Circulation




1,2,5,8 : 9
1,4,7,8 : 10
1,3,6,8 : 4
1,2,6,8 : 1

Elementary Decomposition of Circulation



1,2,5,8 : 9
1,4,7,8 : 10
1,3,6,8 : 4
1,2,6,8 : 1


Elementary Decomposition of Circulation



1,2,5,8 : 9
1,4,7,8 : 10
1,3,6,8 : 4
1,2,6,8 : 1

3,4,7 : 2


Elementary Decomposition of Circulation



1,2,5,8 : 9
1,4,7,8 : 10
1,3,6,8 : 4
1,2,6,8 : 1

3,4,7 : 2

Elementary Decomposition of Circulation



1,2,5,8 : 9
1,4,7,8 : 10
1,3,6,8 : 4
1,2,6,8 : 1
1,4,7,3,6,8 : 4

3,4,7 : 2

Non-negative Transshipment

$D = (V, A)$: a simple directed graph

x : a non-negative edge function

Def. $\text{excess}(v) := x(\delta^{\text{in}}(v)) - x(\delta^{\text{out}}(v))$ is called the x -**excess** of v ; and its additive inverse $\text{deficit}(v)$ is called the x -**deficit** of v ;

Def. A node $v \in V$ is said to be **excessive** (resp, **deficient**, **balanced**) if it has a positive (resp, negative, 0) excess.

Def. If $\text{excess}(v) = b(v)$ for $v \in V$, then x is called a b -TS

Def. The value of a b -TS x : $|x| := (1/2)\|b\|_1$


= total excesses of excessive nodes,

= total deficits of deficient nodes.

Elementary Transshipment

Elementary TS: χ_P along a path or circuit P in A

- $\chi_P(a) = 1$ for each $a \in P$;
- $\chi_P(a) = 0$ for each $a \notin P$.



Elementary Decomposition of TS

$$A^+(x) := \{a \in A : x(a) > 0\}$$

Theorem: There exist a collection \mathcal{P} of **paths** and a collection \mathcal{C} of **circuits** in $A^+(x)$, with positive amounts ε s.t.

- $x = \sum_{P \in \mathcal{P} \cup \mathcal{C}} \varepsilon(P) \chi_P$;
- each path $P \in \mathcal{P}$ is from a deficient node to an excessive node;
- $|\mathcal{P}| + |\mathcal{C}| \leq |A^+(x)| \leq m$;
- moreover, if x is integer-valued, so are all $\varepsilon(P)$.


The decomposition can be found in $O(nm)$ time. Furthermore,

$$|x| = \sum_{P \in \mathcal{P}} \varepsilon(P)$$

Convert to Circulation

If not a circulation,

- add a new vertex v ;
- add an arc (u, v) for each excessive u , put $x(u, v) := x(\delta^{in}(u))$;
- add an arc (v, w) for each deficient w , put $x(v, w) := x(\delta^{out}(w))$.



The new x is a circulation, and a decomposition of the new x gives a decomposition of the original x .

Summary

Single-source shortest paths

- Unit lengths: BFS in $O(m+n)$ time
- Non-negative lengths: Dijkstra's algorithm in $O(m+n\log n)$ time
- No-negative circuits: Bellman-Ford algorithm in $O(mn)$ time, or Dijkstra's algorithm with a given potential

All-to-all shortest paths

- Floyd-Washall algorithm: $O(n^3)$ time
- Johnson's algorithm: $O(n(m+n\log n))$ time.

Minimum mean circuit

- Karp's algorithm: $O(mn)$ time.

Elementary decomposition of circulations and transshipments

Further Topics

k shortest paths

https://en.wikipedia.org/wiki/K_shortest_path_routing

Shortest trail

- Special case of min-cost flow

Shortest path in **Monge DAG**: linear time and space

- Shortest k -link path: $O(\sqrt{k(n-k)}\sqrt{n \log(n-k)})$ time and linear space

Shortest path in **undirected** graphs without negative circuits

- Reduction to **weighted non-bipartite matching**