

Lecture 5. Max-Flow Algorithms

History of Worst-Case Running Times

Year	Discoverer	Method	Asymptotic Time
1951	Dantzig	Simplex	$m n^2 C^\dagger$
1955	Ford, Fulkerson	Augmenting path	$m n C^\dagger$
1970	Edmonds-Karp	Shortest path	$m^2 n$
1970	Edmonds-Karp	Fattest path	$m \log U (m \log n)^\dagger$
1970	Dinitz	Improved shortest path	$m n^2$
1972	Edmonds-Karp, Dinitz	Capacity scaling	$m^2 \log C^\dagger$
1973	Dinitz-Gabow	Improved capacity scaling	$m n \log C^\dagger$
1974	Karzanov	Preflow-push	n^3
1983	Sleator-Tarjan	Dynamic trees	$m n \log n$
1986	Goldberg-Tarjan	FIFO preflow-push	$m n \log (n^2 / m)$
...
2013	Orlin + KTR	Contraction	mn

† Edge capacities are between 1 and C .

Assumptions

Flow network $D = (V, A; c)$


- Simple, bidirected
- c is nonnegative
- Every node is on an s - t path
- No uncapacitated s - t path

Outline

- Augmenting flow by single path-flow
- Augmenting flow by blocking-flow
- Preflow push on arcs

1. Augmenting Flow by Single Path-Flow

Ford-Fulkerson Method





Ford-Fulkerson Method

Ford-Fulkerson (D, s, t, c)

$f \leftarrow 0$


while (there exists f -augmenting path P) $f \leftarrow f \oplus P$

return f




Zwick's flow network

Theorem. The Ford-Fulkerson algorithm may not terminate; moreover, it may converge to a value not equal to the value of the maximum flow.




Valid Ford-Fulkerson sequence


Iteration $4k - 3: r^{2k-1}$




Iteration $4k - 2: r^{2k-1}$



Iteration $4k - 2: r^{2k}$



Iteration $4k: r^{2k}$



Choosing good augmenting paths

- Can find augmenting paths efficiently.
- Few iterations.

Choose augmenting paths with: [Edmonds-Karp 1972, Dinitz 1970]

- Max bottleneck capacity.
- Sufficiently large bottleneck capacity.
- Fewest number of edges.

(Nearly) Widest Augmenting Path




A greedy intuition

[Edmonds -Karp 1970,1972; Dinitz 1973]


Widest augmenting path: max increase, but slow to find

Nearly widest augmenting path: nearly max increase in $O(m)$ time

- Maintain scaling parameter Δ .
- $D_f(\Delta)$: the subgraph of D_f consisting of only arcs with capacity $\geq \Delta$



D_f



$D_f(100)$

Capacity scaling

Assumption. All edge capacities are **integers** with absolute values at most C .

Scaling-Max-Flow (D, s, t, c)

$C \leftarrow \max\{c(a) : a \in A\}$

$\Delta \leftarrow$ smallest power of 2 greater than or equal to C

$f \leftarrow 0$

while ($\Delta \geq 1$)

while (there exists augmenting path P in $D_f(\Delta)$) do $f \leftarrow f \oplus P$

$\Delta \leftarrow \Delta / 2$

return f

The outer while loop (scaling phase) repeats $1 + \lceil \log C \rceil$ times.

Correctness

Invariants.

- At the beginning of Δ -phase, $D_f(2\Delta)$ has no s-t path
- All flow and residual capacity values are integral.

Correctness. When the algorithm terminates, f is a max flow.

Pf.


- By **integrality** invariant, when $\Delta = 1 \Rightarrow D_f(\Delta) = D_f$.
- Upon termination of $\Delta = 1$ phase, there are no augmenting paths. ▀

Optimality gap at the end of each phase

Claim. At the end of a Δ -scaling phase, $val(f^*) - val(f) < m\Delta$.

Pf.

- U : the set of nodes reachable from s in $D_f(\Delta)$
- $val(f^*) - val(f) \leq (\text{residual}) \text{ cut-capacity of } U \text{ in } D_f < m\Delta$



residual network

Running time

Claim. There are at most $2m$ augmentations per scaling phase.

Pf.

- At the beginning of Δ -phase, the optimality gap $< m(2\Delta) = 2m\Delta$.
- Each augmentation in a Δ -phase decreases the gap by at least Δ . ▫

Totally, $O(m \log C)$ augmentations.


Overall running time: $O(m^2 \log C)$ time.

Weakly polynomial, may not terminate for irrational capacities

[Queyranne 1980]


Queyranne's flow network

Theorem. The Edmonds-Karp algorithm may not terminate.




Valid Edmonds-Karp sequence


Iteration $3k - 2$: r^{3k-2}



Iteration $3k - 1$: r^{3k-1}



Iteration $3k$: r^{3k}




Shortest Augmenting Path



Shortest Augmenting Path

[Dinitz 1970, Edmonds-Karp 1972]



Shortest-Augmenting-Path (D, s, t, c)


$f \leftarrow 0$

while (there exists an augmenting path)
 find such a **shortest** such path P using BFS
 $f \leftarrow f \oplus P$

return f



Level graph

- L_f : subgraph of D_f containing of all vertices and edges appearing in some shortest s-t path in D_f
 - Compute in $O(m + n)$ time using BFS by keeping only forward edges (deleting back and side edges).
- P is a shortest s-t path in D_f iff it is an s-t path L_f .





Evolution of level graphs

- Let f and f' be flow before and after a shortest path augmentation.
- Only back edges added to $D_{f'}$, and at least one edge (the bottleneck edge) in D_f (and L_f) is deleted from $D_{f'}$ (and $L_{f'}$)
- Path with back edge has length greater than previous length.



Evolution of level graphs

- The length of the shortest path never decreases.
- If the length of shortest s - t path in $D_{f'}$ does not increase, then the edge set of $L_{f'}$ strictly decreases.



Running time

Phase: successive shortest path augmentations in which the shortest augmenting paths have the same length

Fact: at most n phases, and at most m augmentations per phase

Theorem. The shortest augmenting path algorithm performs at most $O(mn)$ augmentations. The overall running time is $O(m^2n)$.

2. Augmenting Flow by Blocking-Flow




Blocking augmenting flow


Analogy of maximal disjoint paths in the Hopcroft-Karp algorithm for maximum bipartite matching

Def. For a flow f in D , a flow g in L_f is **blocking** if all edges in L_f not saturated by g contain no s-t path

Blocking \nRightarrow Maximum



Maximum \Rightarrow Blocking



Dinitz's Method

Dinitz (D, s, t, c)

$f \leftarrow 0$

while (there exists f -augmenting path)

$L_f \leftarrow$ the level graph of D_f

$g \leftarrow$ a blocking flow in L_f

$f \leftarrow f + g$

return f

Algorithms for finding block-flow:

[Dinitz 1970]: $O(mn)$

[Karzanov 1974]: $O(n^2)$

[Sleator-Tarjan 1983]: $O(m \log n)$




Finding blocking flow via DFS

[Dinitz 1970]



Each iteration starts at s , and acts at current vertex v as follows:

- Case 1: v has a forward edge. Move along a forward edge to the next node.
- Case 2: v has no forward edge.
 - subcase 2.1: $v = s$. Stop
 - subcase 2.2: $v = t$. **Augment** and **delete** bottleneck edges on **path**. If t has no incoming edge, stop; otherwise, move on to the next iteration
 - subcase 2.3: $v \neq s, t$. **Delete** v (and all its incident edges) and move backward to its predecessor.

Demo



Demo



Running time

- Each iteration, at least one **edge** is deleted. There are $O(m)$ iterations.
- Running time of each iteration: $O(n + \text{number of edges deleted})$
- Total running time of finding a blocking flow: $O(mn + m) = O(mn)$
- Total running time of finding maximum flow: $O(n^2 m)$

Finding blocking flow via preflow push/pull

Def. (residual) capacity of a vertex in L_f :



$$\begin{aligned}c(s) &:= c(\delta^{out}(s)) \\c(t) &:= c(\delta^{in}(t)) \\c(v) &:= \min\{c(\delta^{in}(v)), c(\delta^{out}(v))\}\end{aligned}$$

Each iteration:

- Compute a vertex v^* with minimum capacity ε .
- If $\varepsilon > 0$, then
 - Compute an $s - v^*$ flow of value ε by greedy pulling
 - Compute a $v^* - t$ flow of value ε by greedy pushing
- If $c(s)$ or $c(t)$ is ε , return g ;
- **Remove** v^* and its incident edges, update its neighbors' capacity, and repeat.

Greedy pulling/pushing

- Greedy pulling (or pushing) backward (or forward) from v^* to s (or t) **level-by-level**
- At each vertex scan the incoming (or outgoing) edges one at a time:
 - **fully saturate** the edge before going on to the next one.
 - **remove** saturated edges
 - Update the capacities of itself and its neighbors.



Running time

- $O(n)$ iterations: at least one **vertex** is removed in each iteration.
- Running time of each iteration: $O(n + \text{number of edges removed})$
- Total running time of finding blocking flow: $O(n^2 + m) = O(n^2)$
- Total running time of finding maximum flow: $O(n^3)$

3. Preflow Push on Arcs



Overview

- Flow-augmenting approach: maintain a **flow** f and iteratively augment it until no s - t path in A_f (i.e. optimality)
- Preflow push/lift approach: maintain an **s -preflow** f **without** s - t path in A_f and modify it on an arc-by-arc basis until f is a flow, which is optimal
 - Initial s -preflow: the source s saturates all outgoing arcs
 - Subsequently, pick an excessive node $u \neq t$ to discharge its excess towards residual neighbors (including s possibly) "**closer**" to t

Recap: preflow

f : s -preflow

Fact. Every excessive node has at least one residual neighbor and can reach s in the residual graph.

Pf. The elementary decomposition of f has an s - v path

$$P \subseteq A^+(f) \subseteq A_f^{-1}$$

$P^{-1} \subseteq A_f$ is a v - s path in D_f .

Push operation

Condition: $u \neq t$ is excessive and $(u, v) \in A_f$

Push(u, v):

$\varepsilon \leftarrow \min\{f(\delta^{in}(u)), c_f(u, v)\}$; //maximal amount
 $f(u, v) \leftarrow f(u, v) + \varepsilon, f(v, u) \leftarrow -f(u, v)$.

Classification: **balancing** if $\varepsilon = f(\delta^{in}(u))$, **non-balancing** otherwise

- balancing: u becomes balanced;
- non-balancing: (u, v) is removed from A_f


Decision: which u and which residual edge (u, v) ?

Analogy: fluid naturally finds its way "downhill".

Valid heights

$h: V \rightarrow \mathbb{Z}_+$ with $h(s) = n$ (hill middle) and $h(t) = 0$ (hill bottom)

Def. h is **valid** for f if for each $(u, v) \in A_f$ then $h(u) - h(v) \leq 1$



Initial height for initial f : $h(s) = n$ and $h(v) = 0$ for all $v \neq s$.

Properties of valid height

Claim. There exists a valid h for $f \Leftrightarrow$ there is no s - t path in D_f .

Pf. (\Rightarrow) Otherwise, the s - t distance in D_f would be $\geq h(s) - h(t) = n$.

(\Leftarrow) Put all nodes reachable from s at level n , and others at level 0.

Claim. For any valid h and any **excessive** node u , $h(u) < 2n$.

Pf.

$$n > u\text{-}s \text{ distance in } D_f \geq h(u) - h(s) = h(u) - n$$

So, $h(u) < 2n$.

The highest rule and downhill rule

- Choose a **highest** (excessive) u
- If u has a residual neighbor v with $h(v) = h(u) - 1$, push along (u, v)

Claim. $\text{Push}(u, v)$ preserves the validity of h .

Pf. Easy verification at u and also at v .

- Otherwise,

Lift(u): $h(u) := h(u) + 1$.

Claim. $\text{Lift}(u)$ preserves the validity of h .

Treatment on u

```
while ( $u$  is not balanced)
  if ( $\exists (u, v) \in A_f$  with  $h(v) = h(u) - 1$ ) Push( $u, v$ );
  else Lift( $u$ ).
```

The last push is **balancing**, and all others are non-balancing.

Push/Lift Algorithm

```
initialize  $f, h$ ;  
while (there is an excessive node other than  $t$ )  
    pick a highest excessive node  $u \neq t$ ;  
    treat  $u$ ;  
return  $f$ 
```

- The number of lifts is $< 2n^2$
- The number of balancing pushes is $O(n^3)$
- The number of non-balancing push is $O(mn)$
- The total number of operations is $O(n^3)$

Simple $O(n^3)$ -time implementation with linked lists and arrays.

Evolution of heights

- s and t have fixed heights
- Every other node u can never go down and its final height is $< 2n$
 - If u has never been lifted, its final height is 0.
 - Otherwise, it is excessive right after the last lift and hence its final height is $< 2n$.
- The number of lifts per node is $< 2n$
- The total number of lifts is $< 2n^2$.

Number of balancing pushes

Claim: $\leq n - 2$ balancing pushes between any two consecutive lifts,

Pf. Each of them makes one highest excessive node balanced.

The total number of balancing pushes is $O(n^3)$.

Number of non-balancing pushes

- $\forall (u, v) \in A$, a non-balancing **Push**(u, v) can occur at most n times.
- The number of non-balancing pushes is $O(mn)$.

Claim Between two consecutive non-balancing **Push**(u, v), the height of v increases by at least 2.

Pf. Right after the first non-balancing **Push**(u, v), u is above v , and (u, v) is **no longer** a residual edge.

Right before the second non-balancing **Push**(u, v), (u, v) must have become a residual edge again, which must be a consequence of some **Push**(v, u).

However, in order to make **Push**(v, u), we first need for v 's height to increase by at least 2 so that v is above u .

Summary

- Augmenting flow by single path-flow
 - (Nearly) Widest Augmenting Path
 - Shortest Augmenting Path
- Augmenting flow by blocking-flow
 - via DFS
 - via preflow push/pull
- Preflow push on arcs: push/lift
- Still active research on faster **weakly** polynomial-time algorithm
- <https://www.quantamagazine.org/researchers-achieve-absurdly-fast-algorithm-for-network-flow-20220608/>